

mini-project-1

March 20, 2023

0.0.1 Importing Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

0.0.2 Data Collection and Analysis

```
[2]: # Loading data from CSV file to Pandas Dataframe
```

```
df = pd.read_csv('insurance.csv')
```

```
[3]: # First 5 rows of the Dataframe
```

```
df.head()
```

```
[3]:
```

	age	sex	bmi	children	smoker	region	charges	insuranceclaim
0	19	0	27.900	0	1	3	16884.92400	1
1	18	1	33.770	1	0	2	1725.55230	1
2	28	1	33.000	3	0	2	4449.46200	0
3	33	1	22.705	0	0	1	21984.47061	0
4	32	1	28.880	0	0	1	3866.85520	1

```
[4]: # Number of rows and columns
```

```
df.shape
```

```
[4]: (1338, 8)
```

```
[5]: # Getting Information about Dataset
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	age	1338 non-null	int64
1	sex	1338 non-null	int64
2	bmi	1338 non-null	float64
3	children	1338 non-null	int64
4	smoker	1338 non-null	int64
5	region	1338 non-null	int64
6	charges	1338 non-null	float64
7	insuranceclaim	1338 non-null	int64

dtypes: float64(2), int64(6)

memory usage: 83.8 KB

[6]: *# checking for Missing values*

```
df.isnull().sum()
```

```
[6]: age          0
     sex          0
     bmi          0
     children     0
     smoker       0
     region       0
     charges      0
     insuranceclaim 0
     dtype: int64
```

[7]: *# Statical Information of dataset*

```
df.describe()
```

```
[7]:
```

	age	sex	bmi	children	smoker \
count	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	0.505232	30.663397	1.094918	0.204783
std	14.049960	0.500160	6.098187	1.205493	0.403694
min	18.000000	0.000000	15.960000	0.000000	0.000000
25%	27.000000	0.000000	26.296250	0.000000	0.000000
50%	39.000000	1.000000	30.400000	1.000000	0.000000
75%	51.000000	1.000000	34.693750	2.000000	0.000000
max	64.000000	1.000000	53.130000	5.000000	1.000000


	region	charges	insuranceclaim
count	1338.000000	1338.000000	1338.000000
mean	1.515695	13270.422265	0.585202
std	1.104885	12110.011237	0.492871
min	0.000000	1121.873900	0.000000
25%	1.000000	4740.287150	0.000000

50%	2.000000	9382.033000	1.000000
75%	2.000000	16639.912515	1.000000
max	3.000000	63770.428010	1.000000

[19]: *# Checking for Age distribution*

```
plt.figure(figsize=(5,6))
sns.displot(df['age'])
plt.title('Age Distribution')
plt.show()
```

<Figure size 360x432 with 0 Axes>




[20]: *# Checking for BMI distribution*

```
plt.figure(figsize=(5,6))
sns.displot(df['bmi'])
```

```
plt.title('BMI Distribution')  
plt.show()
```

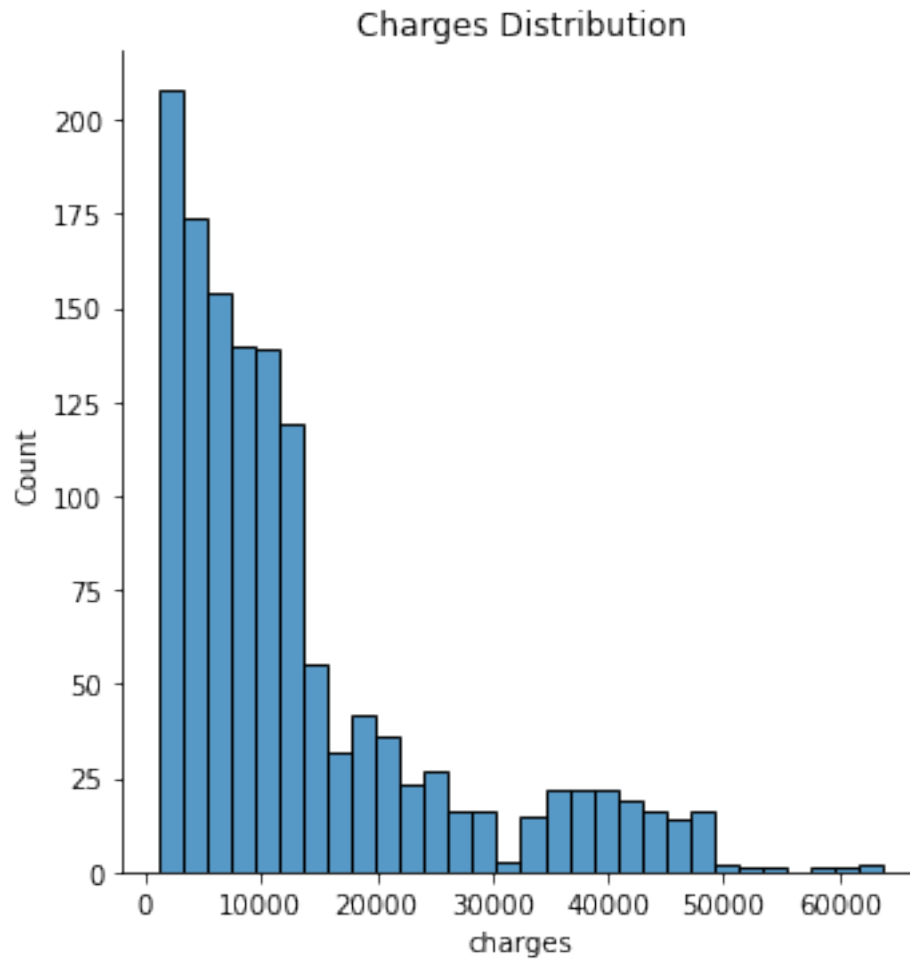
<Figure size 360x432 with 0 Axes>



[21]: *# Checking for Charges distribution*

```
plt.figure(figsize=(5,6))  
sns.displot(df['charges'])  
plt.title('Charges Distribution')  
plt.show()
```

<Figure size 360x432 with 0 Axes>



0.0.3 Splitting Feature & Target

```
[8]: X = df.drop(['insuranceclaim'],axis=1)
     y = df['insuranceclaim']
```

```
[9]: # Feature
```

```
X
```

```
[9]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	3	16884.92400
1	18	1	33.770	1	0	2	1725.55230
2	28	1	33.000	3	0	2	4449.46200
3	33	1	22.705	0	0	1	21984.47061
4	32	1	28.880	0	0	1	3866.85520
...

1333	50	1	30.970	3	0	1	10600.54830
1334	18	0	31.920	0	0	0	2205.98080
1335	18	0	36.850	0	0	2	1629.83350
1336	21	0	25.800	0	0	3	2007.94500
1337	61	0	29.070	0	1	1	29141.36030

[1338 rows x 7 columns]

```
[10]: # Target
```

```
y
```

```
[10]: 0      1
      1      1
      2      0
      3      0
      4      1
      ..
     1333    0
     1334    1
     1335    1
     1336    0
     1337    1
     Name: insuranceclaim, Length: 1338, dtype: int64
```

0.0.4 Splitting data into Training & Testing Data

```
[11]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
      ↪30,random_state=42)
```

```
[12]: # Printing shape
```

```
print("X_train :",X_train.shape)
print("X_test :",X_test.shape)
print("y_train :",y_train.shape)
print("y_test :",y_test.shape)
```

```
X_train : (936, 7)
X_test  : (402, 7)
y_train : (936,)
y_test  : (402,)
```

0.0.5 Model Training

```
[13]: # Loading the Decision Tree Classification model

from sklearn.tree import DecisionTreeClassifier

dt = DecisionTreeClassifier()
```

```
[14]: # Fitting the model

dt.fit(X_train,y_train)
```

```
[14]: DecisionTreeClassifier()
```

0.0.6 Model Evaluation

```
[15]: # Checking Accuracy on Training Data

acc1 = dt.score(X_train,y_train)
print("Accuracy on Training data :",acc1)
```

Accuracy on Training data : 1.0

```
[16]: # Prediction on Testing Data

y_pred = dt.predict(X_test)
print(y_pred)
```

```
[0 1 1 1 1 1 0 1 1 1 0 0 0 1 1 1 0 1 0 1 0 0 1 1 1 1 1 0 1 1 0 1 1 0 0 1 0
 1 1 1 0 1 1 0 0 0 0 0 1 0 1 0 1 1 1 0 0 1 1 1 1 1 0 0 1 0 0 1 0 1 0 0 1 0
 1 1 0 0 0 1 1 1 0 1 1 1 0 1 0 0 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 0 0 1 0 1 1
 1 1 1 1 1 1 1 0 0 1 0 1 1 0 0 1 1 1 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 0 1 0 1
 0 0 1 1 1 1 0 1 1 0 1 1 0 0 1 0 0 1 1 1 1 1 0 1 1 0 1 1 1 1 0 0 1 1 0 1 0
 1 0 1 1 1 0 1 0 1 1 1 0 1 1 0 0 1 1 0 1 1 1 0 1 0 0 0 1 1 0 1 0 1 0 1 1 1
 0 1 1 1 0 1 1 0 1 0 0 1 0 1 0 1 1 1 0 1 1 0 0 1 0 1 0 1 1 0 0 0 1 1 1 1 0
 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 0 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 0 0 0
 1 0 1 1 1 0 0 1 1 1 1 1 0 1 0 0 1 1 0 0 1 1 1 0 1 0 0 1 0 1 0 0 1 1 1 1 1
 1 1 0 0 1 1 1 1 0 0 0 1 0 0 1 0 1 0 0 0 1 1 1 1 1 0 1 1 1 0 1 1 0 1 0 0 0
 1 1 0 0 1 1 1 1 1 1 0 1 0 0 1 1 1 1 1 0 0 1 0 0 1 0 0 1 1 1 1 0]
```

```
[17]: # Checking Accuracy on Testing Data

acc2 = dt.score(X_test,y_test)
print("Accuracy on Tetsing data :",acc2)
```

Accuracy on Tetsing data : 0.9751243781094527

```
[18]: # I used different Classification model on Same dataset and I got the Accuracy ↵  
      ↪as :
```

```
# Logistic regression : 0.815920398  
# Support Vector Machine : 0.6082089552238806  
# Naive bayes score : 0.7723880597014925
```